

The Sapiient Company

AI Capabilities Statement · Human-Response Grounding for AI Interpretability, Control, and Robustness

Who we are

The Sapiient Company builds the cognition layer between machines and humans: a real-time, in-silico read of how a person responds to an AI system. Our core model, Sapiient-1, is a stimulus-to-cortex encoder that predicts whole-cortex fMRI activation (20,484 fsaverage5 vertices at 1 Hz), conditioned per subject. On top of it, our neurosignal layer maps predicted brain-network activity to interpretable, literature-grounded constructs (reward and value, emotional resonance, attention, decision conflict, cognitive load), plus calibrated manipulation and sycophancy indices for auditing AI outputs.

Relevance to the AI Forge strategic thrusts

Interpretability. We add a human-effect axis to interpretability: instead of only explaining what a model represents, we estimate the predicted neural and behavioral impact of its outputs on people, with every construct tied to cited neuroscience and withheld when the signal is confounded [1,2,5,6].

AI control. A calibrated, real-time manipulation and sycophancy signal is a control primitive: it flags when a system is steering an operator through affective or reward drive while suppressing analytic engagement, and can gate or escalate [6,7].

Adversarial robustness. Manipulation is an adversarial attack on the human in the loop, the softest target in a human-AI team. Detecting it protects operator decision quality in contested settings [1,6].

What we could contribute

An open, neural-grounded benchmark for AI-on-human persuasion and manipulation; a shared, auditable construct and metric layer other performers can adopt; and a bridge between the computational-neuroscience and AI-safety communities, which rarely share evaluation tooling today.

Honest scope

Sapiient-1 produces predicted activations, not measured fMRI. Results to date are proof-of-concept scale. We state this plainly in all materials and withhold a construct score when coverage is insufficient or the signal is confounded. All content here is unclassified and non-proprietary.

Contact: Aayush Gandhi, CTO, The Sapiient Company · aayush@thesapiientcompany.com

References

- [1] Knutson, Rick, Wimmer, Prelec & Loewenstein (2007). Neural predictors of purchases. *Neuron* 53:147.
- [2] Bartra, McGuire & Kable (2013). The valuation system: a coordinate-based meta-analysis. *NeuroImage* 76:412.
- [3] Cakir, Cakar, Giriskan & Yurdakul (2018). Neural correlates of purchase behavior through fNIRS. *European Journal of Marketing* 52(1/2):224.
- [4] Kapoor, Sahay, Singh, Pammi & Banerjee (2023). The neural correlates and underlying processes of weak brand choices. *Journal of Business Research* 154:113230.
- [5] Corbetta & Shulman (2002). Control of goal-directed and stimulus-driven attention in the brain. *Nature Reviews Neuroscience* 3:201.
- [6] Botvinick, Braver, Barch, Carter & Cohen (2001). Conflict monitoring and cognitive control. *Psychological Review* 108:624.
- [7] Phan, Wager, Taylor & Liberzon (2002). Functional neuroanatomy of emotion: a meta-analysis. *NeuroImage* 16:331.