


Same model, same content: knowing WHO the person is roughly triples accuracy

Story (22 people, audio + text)



Film (4 people, audio + video)

